

Experiment 2

Linear Regression

Aim

To implement and visualize simple linear regression and multiple linear regression, and to evaluate model performance using standard regression metrics.

Theory

Linear Regression is one of the most fundamental and widely used algorithms in Machine Learning and Statistical Modelling. It belongs to the category of supervised learning algorithms, where the model is trained using labelled data that contains both input variables (features) and the corresponding output variable (target).

The goal of Linear Regression is to identify and quantify the relationship between variables so that future outcomes can be predicted. The algorithm assumes that the relationship between the dependent variable and independent variables can be approximated using a linear equation.

For example, linear regression can be used to model relationships such as House price prediction based on area and number of rooms, Sales prediction based on advertising expenditure, Student performance prediction based on study hours, and Once the model learns the relationship from historical data, it can be used to predict continuous numerical values for new input data.

1. Simple Linear Regression

Simple Linear Regression models the relationship between one independent variable (X) and one dependent variable (Y). The objective is to find the best-fitting straight line, known as the regression line, that represents the relationship between the two variables.

The regression line is determined using the equation:

$$Y = \beta_0 + \beta_1 X + \epsilon$$

This line represents the predicted value of Y for a given value of X.




Figure 1: Linear regression showing slope, intercept, and prediction error

As shown in Figure 1, the relationship between the independent variable X and the dependent variable Y is represented using a linear regression model. The straight line $y = \beta_0 + \beta_1 x$ represents the regression line, which predicts the value of Y for a given value of X . Here, β_0 is the intercept, which indicates the value of Y when $X = 0$, and β_1 represents the slope of the line, showing how much Y changes for a unit change in X . The light blue points in the graph represent the actual observed data points. For a particular value X_i , the model predicts a value of Y on the regression line, known as the predicted value, while the actual data point represents the observed value. The vertical distance between these two values is called the random error

or residual ϵ_i . This residual shows the difference between the predicted value given by the model and the actual observed value. The dashed lines in the graph help visualize the observed value, predicted value, and the slope of the regression line, illustrating how linear regression fits a line that best represents the overall trend of the data.

2. Multiple Linear Regression

Multiple Linear Regression extends the concept of simple linear regression by considering two or more independent variables to predict the dependent variable. This allows the model to capture more complex relationships in the data where multiple factors influence the outcome.

The general equation of multiple linear regression is:

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_n X_n + \epsilon$$

Each coefficient represents the contribution of a specific feature to the predicted value, while keeping other variables constant.

Where,

- Y : Dependent variable (Target)
- X : Independent variable(s) (Features)

- β_0 : Intercept
- $\beta_1, \beta_2, \dots, \beta_n$: Regression coefficients
- ϵ : Error term (residual)




Figure 2: Multiple linear regression showing the regression plane and residual error in a two-feature space

As illustrated in Figure 2, Multiple Linear Regression is a statistical technique used to model the relationship between two or more independent variables and a single dependent variable by fitting a best-fit plane through a three-dimensional scatter plot of data points. The model is expressed as $Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + s$, where Y represents the dependent variable, X_1 and

X_2 are the predictors, β_0 is the intercept, β_1 and β_2 are the coefficients corresponding to each predictor, and s denotes the error term.

The fitted plane represents the predicted values of the dependent variable; however, the actual data points generally do not lie exactly on this plane. The vertical distance between an observed data point and the plane, indicated by the dashed line, is known as the residual or error. The objective of the model is to minimize these residuals, thereby improving prediction accuracy. This approach enables the model to quantify the contribution of each independent variable and helps in understanding how different predictors influence the outcome.

3. Method of Least Squares

The parameters of the regression model (β_0 and β_1) are estimated using the Least Squares Method. The main idea of this method is to determine the regression line that minimizes the total prediction error.

The error is measured using the Residual Sum of Squares (RSS), which is calculated as the sum of squared differences between the observed values and the predicted values.

$$RSS = \sum (Y_i - \hat{Y}_i)^2$$

Where:

- Y_i = actual value
- $\hat{Y}_i$ = predicted value from the regression model

By minimizing this value, the algorithm finds the best possible line that fits the data points.




Figure 3: Illustration of residuals as vertical distances between observed data points and the regression line

Figure 3 illustrates the concept of residuals in linear regression. The X-axis represents the independent variable x , while the Y-axis represents the dependent variable y . The slanted line shown in the figure is the regression line, which represents the predicted relationship between x and y . The scattered colored points represent the actual observed data points. For each data point, a vertical arrow is drawn between the observed value and the regression line. This vertical distance is called the residual, which represents the difference between the actual observed value and the predicted value given by the regression model. Mathematically, the residual is expressed as $residual_i = sample_i - line_i$. The figure also shows the formula for the Sum of Squared Residuals (SSR), $\sum (residual_k)^2$, which is used in linear regression to measure how well the regression line fits the data. A smaller sum of squared residuals indicates a better fit of the model to the observed data. The point labeled $C(x1, y1)$ demonstrate a specific example where the vertical difference between the observed point and the predicted point on the regression line represents the residual.

4. Interpretation of Regression Coefficients

The coefficients in a regression model provide meaningful insights into the relationship between variables.

- **Intercept (β_0):** Represents the value of the dependent variable when all independent variables are zero.
- **Slope (β_1):** Indicates how much the dependent variable changes when the independent variable increases by one unit.
- **Regression coefficients ($\beta_1, \beta_2, \dots$):** Show the strength and direction of influence of each feature.

For example, if $\beta_1 = 2$, it means that for every unit increase in X , the value of Y increases by 2 units.

5. Algorithm

- **Step 1:** Let X be the input features and (Y) be output values
- **Step 2:** Assume a linear relationship: $y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n$
 - β_0 is the intercept (value when all $x = 0$)
 - $\beta_1, \beta_2, \dots$ are coefficients (weights) for each feature
- **Step 3:** Define the error (residual) for each data point:
 - Error = Actual value - Predicted value
 - $e_i = y_i - (\beta_0 + \beta_1 x_{1i} + \beta_2 x_{2i} + \dots)$
- **Step 4:** Calculate total error using Sum of Squared Errors (SSE):
 - $SSE = \sum (e_i)^2 = \sum (y_i - \hat{y}_i)^2$
- **Step 5:** Find coefficients that minimize SSE using Normal Equation:
 - $\beta = (X^T X)^{-1} X^T y$
 - Where X is the feature matrix and Y is the target vector
- **Step 6:** For new data, predict using:
 - $\hat{y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n$

6. Assumptions of Linear Regression

1. **Linearity** – The relationship between independent and dependent variables is linear.
2. **Independence** – Observations are independent of each other.
3. **Homoscedasticity** – Constant variance of residuals across all values of X .
4. **Normality** – Residuals are normally distributed.

7. Applications of Linear Regression

Linear Regression is widely used in many real-world applications, including:

- Sales forecasting
- Price prediction
- Risk assessment
- Demand estimation
- Economic trend analysis
- Healthcare data analysis

Due to its simplicity and interpretability, linear regression is often used as a baseline model in many machine learning problems.

8. Merits of Linear Regression

- **Simplicity and Interpretability:** Linear Regression is easy to understand and implement. Model coefficients have clear statistical meaning, indicating the magnitude and direction of feature influence on the target variable.
- **Effective for Linearly Related Data:** Performs well when the true relationship between features and target is approximately linear and assumptions are reasonably satisfied.

9. Demerits of Linear Regression

- **Linearity Assumption:** The model cannot capture non-linear relationships unless manual feature transformations (e.g., polynomial terms) are applied.
- **Sensitivity to Outliers:** Least Squares estimation squares residuals, giving excessive weight to outliers, which can significantly distort the model.
- **Multicollinearity Issues:** High correlation among independent variables leads to unstable coefficient estimates and reduced interpretability in Multiple Linear Regression.

Procedure

The Objective of this experiment is to implement Simple Linear Regression on Salary Dataset to regress *Salaries* with corresponding *Years of Experience*. Then to implement Multiple Linear Regression on Car Dataset using categorical features – Fuel types (Petrol, Diesel or CNG), Seller types (Dealer or Individual) and Transmission (Manual or Automatic) – to regress on Selling Price of cars.

Part 1: Simple Linear Regression

1. Load 'Salary_dataset.csv'
2. Define Years of Experience as X and Salary as Y.

3. Split the dataset into training and testing sets.
4. Train the Linear Regression model.
5. Evaluate performance using MAE, MSE, RMSE, and R^2 .
6. Plot the:
 - Scatter plot of actual data
 - Regression line
 - Residual distribution

Part 2: Multiple Linear Regression

1. Import required libraries: pandas, numpy, matplotlib, seaborn, and sklearn.
2. Load the dataset 'car data.csv'
3. Perform exploratory data analysis using:
 - head(), tail(), info(), describe()
4. Check missing values using isnull().sum().
5. Encode categorical variables:
 - Fuel_Type: Petrol = 0, Diesel = 1, CNG = 2
 - Seller_Type: Dealer = 0, Individual = 1
 - Transmission: Manual = 0, Automatic = 1
6. Define features X by removing Car_Name and Selling_Price.
7. Define target variable Y as Selling_Price.
8. Split the dataset into training and testing sets.
9. Train the Linear Regression model.
10. Evaluate the model using MAE, MSE, RMSE, and R^2 .
Plot residuals vs predicted values.

Linear Regression - Practical Implementation

Option A: Single Variable Regression

Step 1: Importing Libraries

Importing Libraries

```
import pandas as pd
import numpy as np
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
print("Libraries Imported");
```

Output:

```
Libraries Imported
```

Step 2: Loading Dataset

Read and store the salary dataset

```
data = pd.read_csv('Salary_dataset.csv')
print("Dataset reading complete")
```

Output:

```
Dataset reading complete
```

Step 3: Data Analysis

Display the first 5 rows of the dataset

```
data.head()
```

Output:

index	YearsExperience	Salary
0	1.2	39344
1	1.4	46206
2	1.6	37732
3	2.1	43526
4	2.3	39892

Display the last 5 rows of the dataset

```
data.tail()
```

Output:

index	YearsExperience	Salary
25	9.1	105583
26	9.6	116970
27	9.7	112636
28	10.4	122392
29	10.6	121873

Summary statistics for numerical columns

```
data.describe()
```

Output:

	YearsExperience	Salary
count	30.000000	30.000000
mean	5.413333	76004.000000
std	2.837888	27414.429785
min	1.200000	37732.000000
25%	3.300000	56721.750000
50%	4.800000	65238.000000
75%	7.800000	100545.750000
max	10.600000	122392.000000

Concise summary of a DataFrame

```
data.info()
```

Output:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30 entries, 0 to 29
Data columns (total 2 columns):
#   Column          Non-Null Count  Dtype
---  -
0   YearsExperience  30 non-null    float64
1   Salary          30 non-null    int64
dtypes: float64(1), int64(1)
memory usage: 612.0 bytes
```

Detect missing values

```
data.isnull().sum()
```

Output:

```
YearsExperience    0
Salary            0
dtype: int64
```

Return a tuple representing the dimensionality of the DataFrame

```
data.shape
```

Output:
(30, 2)

Step 4: Data Preprocessing

Convert YearsExperience to the nearest lower integer by removing decimal values

```
data['YearsExperience'] = np.floor(data['YearsExperience']).astype(int)
data
```

Output:
Dataset after Preprocessing:

Output:

index	YearsExperience	Salary
0	1	39344
1	1	46206
2	1	37732
3	2	43526
4	2	39892
5	3	56643
6	3	60151
7	3	54446
8	3	64446
9	3	57190
10	4	63219
11	4	55795
12	4	56958
13	4	57082
14	4	61112

... 15 more rows (total: 30 rows)

Step 5: Model Training

Create feature matrix X and target vector Y

```
X = data.drop(['Salary'], axis=1)
print("Feature matrix X created.")
```

Output:
Feature matrix X created.

Display the first 5 rows of the feature matrix X

```
X.head()
```

Output:

	YearsExperience
0	1
1	1
2	1
3	2
4	2

Create target vector Y

```
Y = data['Salary']
print("Target vector Y created.")
```

Output:

```
Target vector Y created.
```

Display the first 5 rows of the target vector Y

```
Y.head()
```

Output:


```
0    39344
1    46206
2    37732
3    43526
4    39892
Name: Salary, dtype: int64
```

Split the dataset into training and testing sets

```
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.15, random_state =
123)
model = LinearRegression()
model.fit(X_train, Y_train)
print("Model Trained")
```

Output:

```
Model Trained
```



Step 6: Model Evaluation

Generate predictions for training and testing datasets using the trained model

```
Y_train_pred = model.predict(X_train)
Y_test_pred = model.predict(X_test)
print("Predictions generated for both Train and Test sets.")
```

Output:

```
Predictions generated for both Train and Test sets.
```

Evaluate model performance on training data using MAE, MSE, RMSE, and R2 score

```
mae = mean_absolute_error(Y_train, Y_train_pred)
mse = mean_squared_error(Y_train, Y_train_pred)
rmse = np.sqrt(mse)
r2 = r2_score(Y_train, Y_train_pred)

print("Mean Absolute Error (MAE):", mae)
print("Mean Squared Error (MSE):", mse)
print("Root MSE (RMSE):", rmse)
print("R2 Score:", r2)
```

Output:

```
Mean Absolute Error (MAE): 5457.139011764704
Mean Squared Error (MSE): 39230345.16254117
Root MSE (RMSE): 6263.413219845963
R2 Score: 0.9422825046716273
```

Evaluate model performance on testing data using MAE, MSE, RMSE, and R2 score

```
mae = mean_absolute_error(Y_test, Y_test_pred)
mse = mean_squared_error(Y_test, Y_test_pred)
rmse = np.sqrt(mse)
r2 = r2_score(Y_test, Y_test_pred)

print("Mean Absolute Error (MAE):", mae)
print("Mean Squared Error (MSE):", mse)
print("Root MSE (RMSE):", rmse)
print("R2 Score:", r2)
```

Output:

```
Mean Absolute Error (MAE): 3245.035764705882
Mean Squared Error (MSE): 20356724.219709344
Root MSE (RMSE): 4511.842663447978
R2 Score: 0.9774778470484203
```

Retrieve the learned coefficient (slope) of the linear regression model

```
model.coef_
```

Output:

```
array([9704.95294118])
```

Retrieve the intercept (bias term) of the linear regression model

```
model.intercept_
```

Output:


```
np.float64(26104.915294117643)
```

Plot actual training data points and the fitted linear regression line for visualization

```
plt.scatter(X_train, Y_train, label='Actual Salary Data')
plt.plot(X_train, Y_train_pred, color='red', label='Fitted Regression Line')
plt.xlabel('Years of Experience'); plt.ylabel('Salary')
plt.title('Linear Regression Fit')
plt.legend(); plt.show()
```

Output:


```
Linear Regression Fit (Training Data)
```

**# Plot actual testing data points and the fitted linear regression line to evaluate model performance**

```
plt.scatter(X_test, Y_test, label='Actual Salary Data')
plt.plot(X_test, Y_test_pred, color='red', label='Fitted Regression Line')
plt.xlabel('Years of Experience'); plt.ylabel('Salary')
plt.title('Linear Regression Fit')
plt.legend(); plt.show()
```

Output:

```
Linear Regression Fit (Testing Data)
```




Plot residuals against predicted values on train data to check error patterns and model assumptions

```
residuals = Y_train - Y_train_pred
plt.scatter(Y_train_pred, residuals); plt.axhline(y=0, color='red', linestyle='--')
plt.xlabel("Predicted Salary")
plt.ylabel("Error (Residual)")
plt.title("Predicted vs Residuals on Train Set"); plt.show()
```

Output:

Predicted vs Residuals (Train Set)




Plot residuals versus predicted values on test data to assess model generalization and error distribution

```
residuals = Y_test - Y_test_pred
plt.scatter(Y_test_pred, residuals)
plt.axhline(y=0, color='red', linestyle='--') # reference line for zero error
plt.xlabel("Predicted Salary")
plt.ylabel("Error (Residual)")
plt.title("Predicted vs Residuals on Test Set")
plt.show()
```

Output:

Predicted vs Residuals (Test Set)

Predicted vs Residuals on Test Set



Predict test values, compare actual vs predicted results, and display the top 5 most accurate predictions

```
Y_test_pred = model.predict(X_test)

comparison_df = pd.DataFrame({
    "Actual Salary": Y_test,
    "Predicted Salary": np.round(Y_test_pred, 2),
    "Error": np.round(Y_test - Y_test_pred, 2),
    "Absolute Error": np.round(abs(Y_test - Y_test_pred), 2)
})

best_predictions = comparison_df.sort_values(by="Absolute Error").head(5)
print(best_predictions)
```

Output:

Best Predictions:

Output:

	Actual Selling Price	Predicted Selling Price	Error	Absolute Error
7	54446	55219.77	-773.77	773.77
29	121873	123154.44	-1281.44	1281.44
5	56643	55219.77	1423.23	1423.23
26	116970	113449.49	3520.51	3520.51
8	64446	55219.77	9226.23	9226.23

Option B: Multi-Variable Regression

Step 1: Importing Libraries

Importing Libraries

```
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
print("Libraries Imported")
```

Output:

```
Libraries Imported
```

Step 2: Loading Dataset

Read and store the Car dataset

```
data = pd.read_csv('car_data.csv')
print("Dataset reading complete")
```

Output:

```
Dataset reading complete
```

Step 3: Data Analysis

Display specific selected rows from the dataset using their index positions

```
rows_to_show = [31, 144, 42, 114, 89]
data.loc[rows_to_show]
```

Output:

	Car_Name	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner
31	ritz	2011	2.35	4.89	54200	Petrol	Dealer	Manual	0
144	Bajaj Pulsar NS 200	2014	0.60	0.99	25000	Petrol	Individual	Manual	0
42	sx4	2008	1.95	7.15	58000	Petrol	Dealer	Manual	0
114	Royal Enfield Classic 350	2015	1.15	1.47	17000	Petrol	Individual	Manual	0
89	etios g	2014	4.75	6.76	40000	Petrol	Dealer	Manual	0

Display the first 5 rows of the dataset

```
data.head()
```

Output:

	Car_Name	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner
0	ritz	2014	3.35	5.59	27000	Petrol	Dealer	Manual	0
1	sx4	2013	4.75	9.54	43000	Diesel	Dealer	Manual	0
2	ciaz	2017	7.25	9.85	6900	Petrol	Dealer	Manual	0
3	wagon r	2011	2.85	4.15	5200	Petrol	Dealer	Manual	0
4	swift	2014	4.60	6.87	42450	Diesel	Dealer	Manual	0

Display the last 5 rows of the dataset

```
data.tail()
```

Output:

	Car_Name	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner
296	city	2016	9.50	11.6	33988	Diesel	Dealer	Manual	0
297	brio	2015	4.00	5.9	60000	Petrol	Dealer	Manual	0
298	city	2009	3.35	11.0	87934	Petrol	Dealer	Manual	0
299	city	2017	11.50	12.5	9000	Diesel	Dealer	Manual	0
300	brio	2016	5.30	5.9	5464	Petrol	Dealer	Manual	0

Show statistical summary of numerical columns

```
data.describe()
```

Output:

	Year	Selling_Price	Present_Price	Kms_Driven	Owner
count	301.000000	301.000000	301.000000	301.000000	301.000000
mean	2013.627907	4.661296	7.628472	36947.205980	0.043189
std	2.891554	5.082812	8.644115	38886.883882	0.247915
min	2003.000000	0.100000	0.320000	500.000000	0.000000
25%	2012.000000	0.900000	1.200000	15000.000000	0.000000
50%	2014.000000	3.600000	6.400000	32000.000000	0.000000
75%	2016.000000	6.000000	9.900000	48767.000000	0.000000
max	2018.000000	35.000000	92.600000	500000.000000	3.000000

Display dataset structure and data types

```
data.info()
```

Output:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 301 entries, 0 to 300
Data columns (total 9 columns):
```

Output:

#	Column	Non-Null Count	Dtype
0	Car_Name	301 non-null	object
1	Year	301 non-null	int64
2	Selling_Price	301 non-null	float64
3	Present_Price	301 non-null	float64
4	Kms_Driven	301 non-null	int64
5	Fuel_Type	301 non-null	object
6	Seller_Type	301 non-null	object
7	Transmission	301 non-null	object
8	Owner	301 non-null	int64

Output:

```
dtypes: float64(2), int64(3), object(4)
memory usage: 21.3+ KB
```

Check the number of missing values in each column

```
data.isnull().sum()
```

Output:

Column	0
Car_Name	0
Year	0
Selling_Price	0
Present_Price	0
Kms_Driven	0
Fuel_Type	0
Seller_Type	0
Transmission	0
Owner	0

Output:

```
dtype: int64
```

Show the number of rows and columns in the dataset

```
data.shape
```

Output:

```
(301, 9)
```

Count occurrences of each Fuel_Type in the dataset

```
data['Fuel_Type'].value_counts()
```

Output:

Fuel_Type	count
Petrol	239
Diesel	60
CNG	2

Output:

dtype: int64

Count occurrences of each Seller_Type in the dataset

```
data['Seller_Type'].value_counts()
```

Output:

Seller_Type	count
Dealer	195
Individual	106

Output:

dtype: int64

Count occurrences of each Transmission type in the dataset

```
data['Transmission'].value_counts()
```

Output:

Transmission	count
Manual	261
Automatic	40

Output:

dtype: int64

Step 4: Data Preprocessing

Encode Fuel_Type column by mapping categorical values to numerical (Petrol:0, Diesel:1, CNG:2)

```
data = data.replace({'Fuel_Type':{'Petrol':0, 'Diesel':1, 'CNG':2}})
data
```

Output:

Dataset after Fuel_Type Encoding:

Output:

	Car_Name	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner
0	ritz	2014	3.35	5.59	27000	0	Dealer	Manual	0
1	sx4	2013	4.75	9.54	43000	1	Dealer	Manual	0
2	ciaz	2017	7.25	9.85	6900	0	Dealer	Manual	0
3	wagon r	2011	2.85	4.15	5200	0	Dealer	Manual	0
4	swift	2014	4.60	6.87	42450	1	Dealer	Manual	0

Linear Regression - Practical Implementation

...	...	...	...	...	...	...	...	...	...
296	city	2016	9.50	11.60	33988	1	Dealer	Manual	0
297	brio	2015	4.00	5.90	60000	0	Dealer	Manual	0
298	city	2009	3.35	11.00	87934	0	Dealer	Manual	0
299	city	2017	11.50	12.50	9000	1	Dealer	Manual	0
300	brio	2016	5.30	5.90	5464	0	Dealer	Manual	0

Output:

301 rows x 9 columns

Encode Seller_Type column by mapping categorical values to numerical (Dealer:0, Individual:1)

```
data = data.replace({'Seller_Type':{'Dealer':0, 'Individual':1}})
data
```

Output:

Dataset after Seller_Type Encoding:

Output:

	Car_Name	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner
0	ritz	2014	3.35	5.59	27000	0	0	Manual	0
1	sx4	2013	4.75	9.54	43000	1	0	Manual	0
2	ciaz	2017	7.25	9.85	6900	0	0	Manual	0
3	wagon r	2011	2.85	4.15	5200	0	0	Manual	0
4	swift	2014	4.60	6.87	42450	1	0	Manual	0
...	...	...	...	...	...	...	...	...	...
296	city	2016	9.50	11.60	33988	1	0	Manual	0
297	brio	2015	4.00	5.90	60000	0	0	Manual	0
298	city	2009	3.35	11.00	87934	0	0	Manual	0
299	city	2017	11.50	12.50	9000	1	0	Manual	0
300	brio	2016	5.30	5.90	5464	0	0	Manual	0

Output:

301 rows x 9 columns

Encode Transmission column by mapping categorical values to numerical (Manual:0, Automatic:1)

```
data = data.replace({'Transmission':{'Manual':0, 'Automatic':1}})
data
```

Output:

Dataset after Transmission Encoding:

Output:

	Car_Name	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner
0	ritz	2014	3.35	5.59	27000	0	0	0	0
1	sx4	2013	4.75	9.54	43000	1	0	0	0
2	ciaz	2017	7.25	9.85	6900	0	0	0	0
3	wagon r	2011	2.85	4.15	5200	0	0	0	0
4	swift	2014	4.60	6.87	42450	1	0	0	0
...	...	...	...	...	...	...	...	...	...
296	city	2016	9.50	11.60	33988	1	0	0	0

297	brio	2015	4.00	5.90	60000	0	0	0	0
298	city	2009	3.35	11.00	87934	0	0	0	0
299	city	2017	11.50	12.50	9000	1	0	0	0
300	brio	2016	5.30	5.90	5464	0	0	0	0

Output:

301 rows x 9 columns

Step 5: Model Training

Create feature matrix X by dropping Car_Name and Selling_Price columns

```
X = data.drop(['Car_Name', 'Selling_Price'], axis=1)
print("Feature matrix X created by dropping 'Car_Name' and 'Selling_Price' columns.")
```

Output:

Feature matrix X created by dropping 'Car_Name' and 'Selling_Price' columns.

Create target vector Y containing the Selling_Price column

```
Y = data['Selling_Price']
print("Target vector Y created from 'Selling_Price' column.")
```

Output:

Target vector Y created from 'Selling_Price' column.

Display the first 5 rows of the feature matrix X

```
X.head()
```

Output:

Feature Matrix (X) Preview:

Output:

	Year	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner
0	2014	5.59	27000	0	0	0	0
1	2013	9.54	43000	1	0	0	0
2	2017	9.85	6900	0	0	0	0
3	2011	4.15	5200	0	0	0	0
4	2014	6.87	42450	1	0	0	0

Display the first 5 values of the target vector Y

```
Y.head()
```

Output:

Target Vector (Y) Preview:

Output:

	Selling_Price
0	3.35
1	4.75
2	7.25
3	2.85
4	4.60

Output:


dtype: float64

Split the dataset into training and testing sets, then train a Linear Regression model

```
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.15, random_state = 123)
model = LinearRegression()
model.fit(X_train, Y_train)
print("Model Trained Successfully")
```

Output:

Model Trained Successfully



Step 6: Model Evaluation

Generate predictions for training and testing datasets using the trained model

```
Y_train_pred = model.predict(X_train)
Y_test_pred = model.predict(X_test)
print("Predictions generated for both Train and Test sets.")
```

Output:

Predictions generated for both Train and Test sets.

Evaluate model performance on training data using MAE, MSE, RMSE, and R2 score

```
mae = mean_absolute_error(Y_train, Y_train_pred)
mse = mean_squared_error(Y_train, Y_train_pred)
rmse = np.sqrt(mse)
r2 = r2_score(Y_train, Y_train_pred)

print("Mean Absolute Error (MAE):", mae)
print("Mean Squared Error (MSE):", mse)
print("Root MSE (RMSE):", rmse)
print("R^2 Score:", r2)
```

Output:

```
Mean Absolute Error (MAE): 1.1605491562123338
Mean Squared Error (MSE): 3.015625977191586
Root MSE (RMSE): 1.7365557800403608
R^2 Score: 0.8869753499147778
```

Evaluate model performance on testing data using MAE, MSE, RMSE, and R2 score

```
mae = mean_absolute_error(Y_test, Y_test_pred)
mse = mean_squared_error(Y_test, Y_test_pred)
rmse = np.sqrt(mse)
r2 = r2_score(Y_test, Y_test_pred)

print("Mean Absolute Error (MAE):", mae)
print("Mean Squared Error (MSE):", mse)
print("Root MSE (RMSE):", rmse)
print("R^2 Score:", r2)
```

Output:

```
Mean Absolute Error (MAE): 1.3160017493663516
Mean Squared Error (MSE): 3.8164774430610064
Root MSE (RMSE): 1.953580672684903
R^2 Score: 0.8114116982277362
```

Retrieve the learned coefficients and intercept of the linear regression model

```
print("Coefficients:", model.coef_)
print("Intercept:", model.intercept_)
```

Output:

```
Coefficients: [ 3.92825483e-01  4.39836584e-01 -5.25747442e-06  1.38880051e+00 -1.25197058e+00
 1.44799328e+00 -7.41998907e-01]
Intercept: -789.4879877881141
```


Plot residuals against predicted values on training data to analyze error patterns

```
residuals = Y_train - Y_train_pred
plt.scatter(Y_train_pred, residuals)
plt.axhline(0, color='red', linestyle='--', linewidth=2)
plt.xlabel("Predicted Price")
plt.ylabel("Residual (Actual - Predicted)")
plt.title("Residual vs Predicted Plot on Train Set")
plt.show()
```

Output:

```
Residual vs Predicted Plot on Train Set
```


Residual vs Predicted Plot on Train Set




Plot residuals against predicted values on testing data to assess model generalization

```
residuals = Y_test - Y_test_pred
plt.scatter(Y_test_pred, residuals)
plt.axhline(0, color='red', linestyle='--', linewidth=2)
plt.xlabel("Predicted Price")
plt.ylabel("Residual (Actual - Predicted)")
plt.title("Residual vs Predicted Plot on Test Set")
plt.show()
```

Output:

Residual vs Predicted Plot on Test Set

Residual vs Predicted Plot on Test Set



Choose one random test sample from the dataset to compare the actual and predicted selling prices.

```
test_samples = data.sample(n=5, random_state=42)
display(test_samples)
print("Interactive Prediction Table Loaded")
```

Output:

```
Interactive Prediction Table Loaded
Choose one random test sample.
```

Output:

	Car_Name	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner
--	----------	------	---------------	---------------	------------	-----------	-------------	--------------	-------

Interactive Features (Static View)

Interactive Feature: Test Sample Predictions

Clicking a row in the test samples table shows the model prediction. Below are all 5 test samples with their prediction results:

Index	Car Name	Year	Actual	Present	Kms	Fuel	Seller	Trans	Owner	Predicted	Error
31	ritz	2011	2.35	4.89	54200	Petrol	Dealer	Manual	0	2.35	0.00
144	Bajaj Pulsar NS 200	2014	0.60	0.99	25000	Petrol	Individual	Manual	0	0.71	-0.11
42	sx4	2008	1.95	7.15	58000	Petrol	Dealer	Manual	0	2.15	-0.20
114	Royal Enfield 350	2015	1.15	1.47	17000	Petrol	Individual	Manual	0	1.36	-0.21
89	etios g	2014	4.75	6.76	40000	Petrol	Dealer	Manual	0	4.43	0.32